

APACHE SPARK


DS8003 – MGT OF BIG DATA AND TOOLS

Ryerson University

Instructor: Kanchana Padmanabhan

Big Data History

2



Big Data – New Trend

3

Compare Search terms ▼

Apache Hadoop

Search term

Apache Spark

Search term


Apache Flink

Search term

+ Add term

Interest over time ?

☐ News headlines ☐ Forecast ?



New TeraSort World Record!

4

- The previous world record was 72 minutes, set by Yahoo using a Hadoop MapReduce cluster of 2100 nodes.
- New world record set by Spark using Spark on 206 EC2 nodes, basically with **3X** faster using **10X** fewer machines.

	Hadoop MR Record	Spark Record	Spark 1 PB
Data Size	102.5 TB	100 TB	1000 TB
Elapsed Time	72 mins	23 mins	234 mins
# Nodes	2100	206	190
# Cores	50400 physical	6592 virtualized	6080 virtualized
Cluster disk throughput	3150 GB/s (est.)	618 GB/s	570 GB/s
Sort Benchmark Daytona Rules	Yes	Yes	No
Network	dedicated data center, 10Gbps	virtualized (EC2) 10Gbps network	virtualized (EC2) 10Gbps network
Sort rate	1.42 TB/min	4.27 TB/min	4.27 TB/min
Sort rate/node	0.67 GB/min	20.7 GB/min	22.5 GB/min

Contributors

Commits

Code frequency

Punch card

Network

Members

Mar 28, 2010 – Mar 30, 2015

Contributions to master, excluding merge commits


Contributions: **Commits** ▾



mateiz

823 commits / 1,627,101 ++ / 1,150,457 --

#1



pwendell

622 commits / 30,046 ++ / 21,712 --

#2



Spark Adoptions

6



- Yahoo! – Personalization and ad analytics
- Conviva – Real-time video stream optimization
- Ooyala – Cross-device personalized video experience
- Groupon, Shopify, Alibaba, Taobao, Tencent, etc...

Spark Hall of Fame

7

Spark “Hall of Fame”



LARGEST CLUSTER

Tencent
(8000+ nodes)



LARGEST SINGLE-DAY INTAKE

Tencent
(1PB+ /day)



LONGEST-RUNNING JOB

Alibaba
(1 week on 1PB+ data)



LARGEST SHUFFLE

Databricks PB Sort
(1PB)



MOST INTERESTING APP

Jeremy Freeman
Mapping the Brain at Scale
(with lasers!)

Apache Spark History

8

- ❑ Spark was initially started by Matei Zaharia at UC Berkeley AMPLab in 2009
- ❑ Open sourced in 2010
- ❑ Donated to Apache Foundation in 2013
- ❑ Became an Apache Top-Level Project in Feb 2014



The Berkeley AMPLab

9

Governmental and industrial funding:



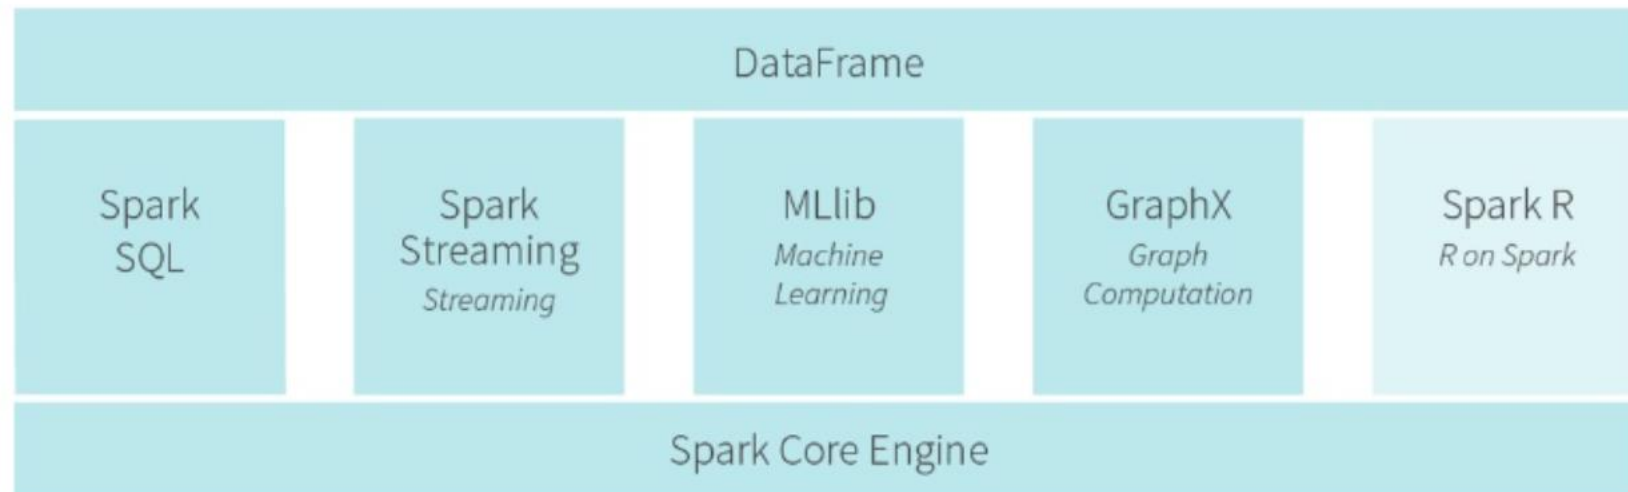
GE imagination at work



Goal: Next generation of open source data analytics stack for industry & academia:
Berkeley Data Analytics Stack (BDAS)

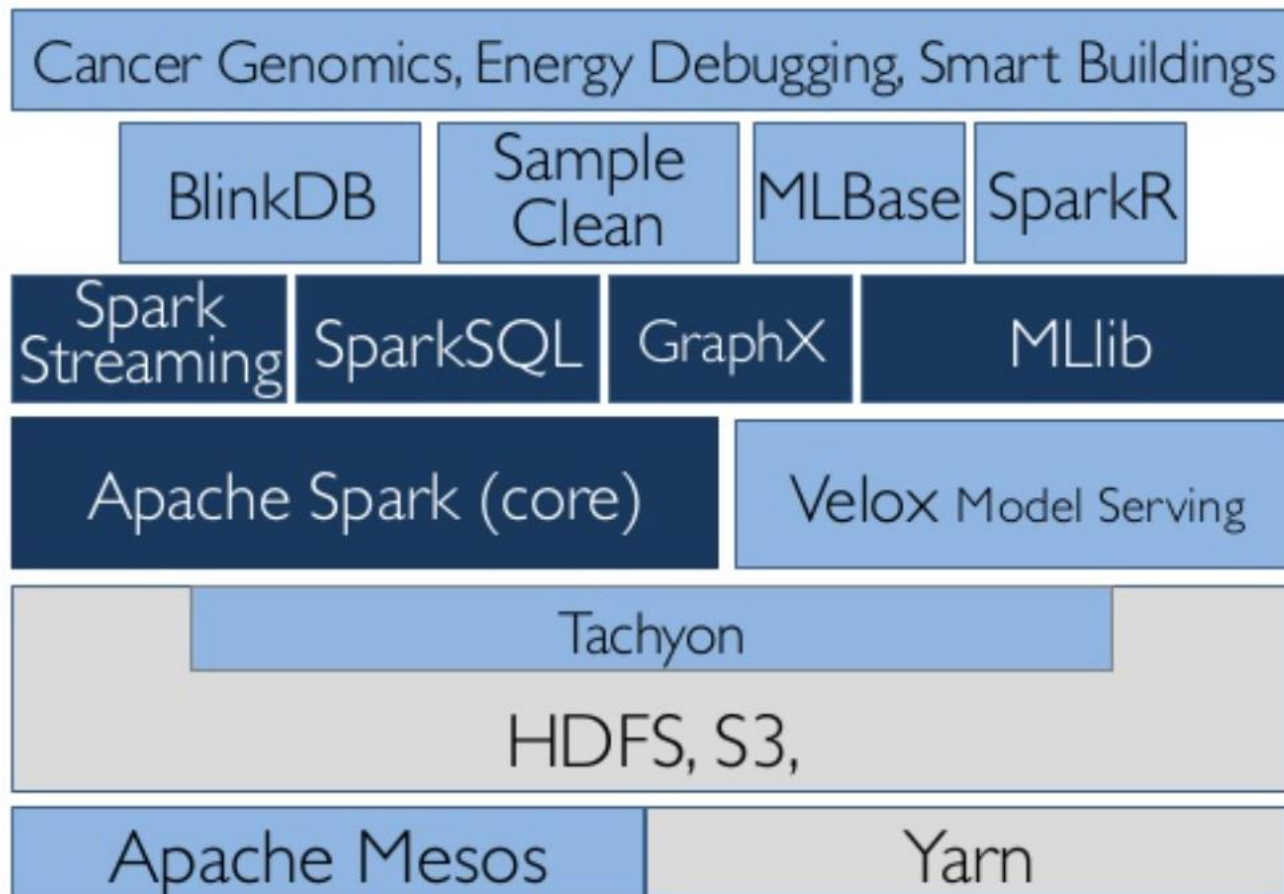
- Databricks is a company founded by the creators of Apache Spark, that aims to help clients with cloud-based big data processing using Spark

Alpha / Pre-alpha



The BDAS – Berkeley Data Analytics Stack

11




Unified Data Platform

12

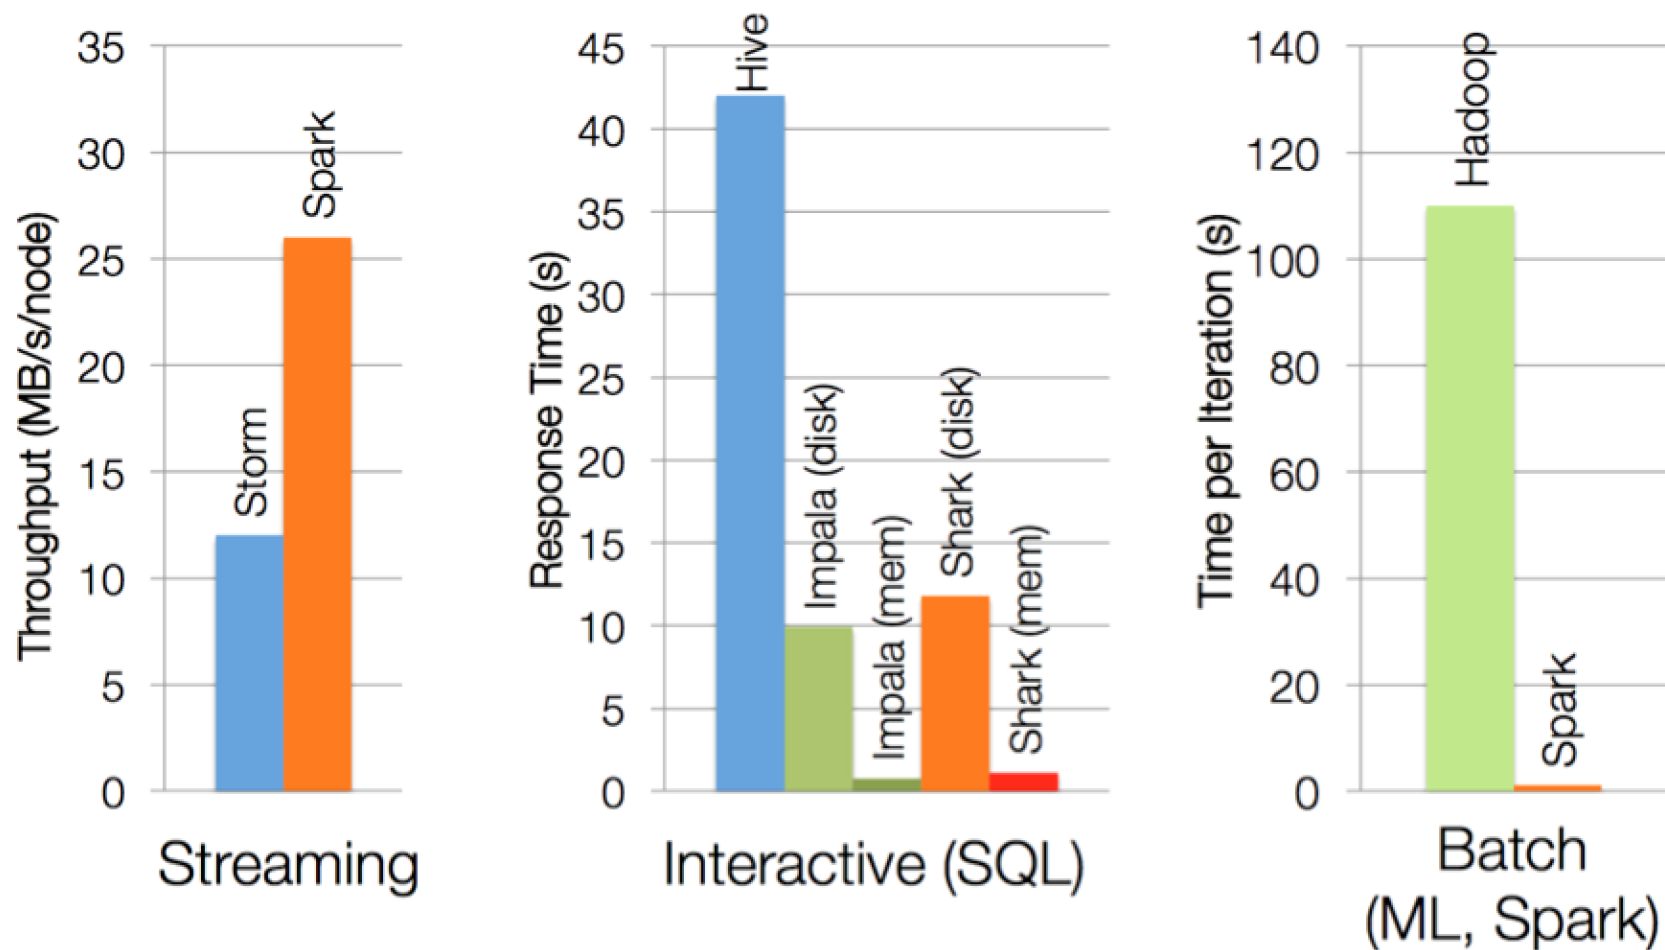
- A unified platform that supports many data processing needs including
 - ▣ Batch processing (Spark)
 - ▣ Stream processing (SparkX)
 - ▣ Interactive (Spark SQL)
 - ▣ Iterative (MLlib, GraphX)

One size fits many!



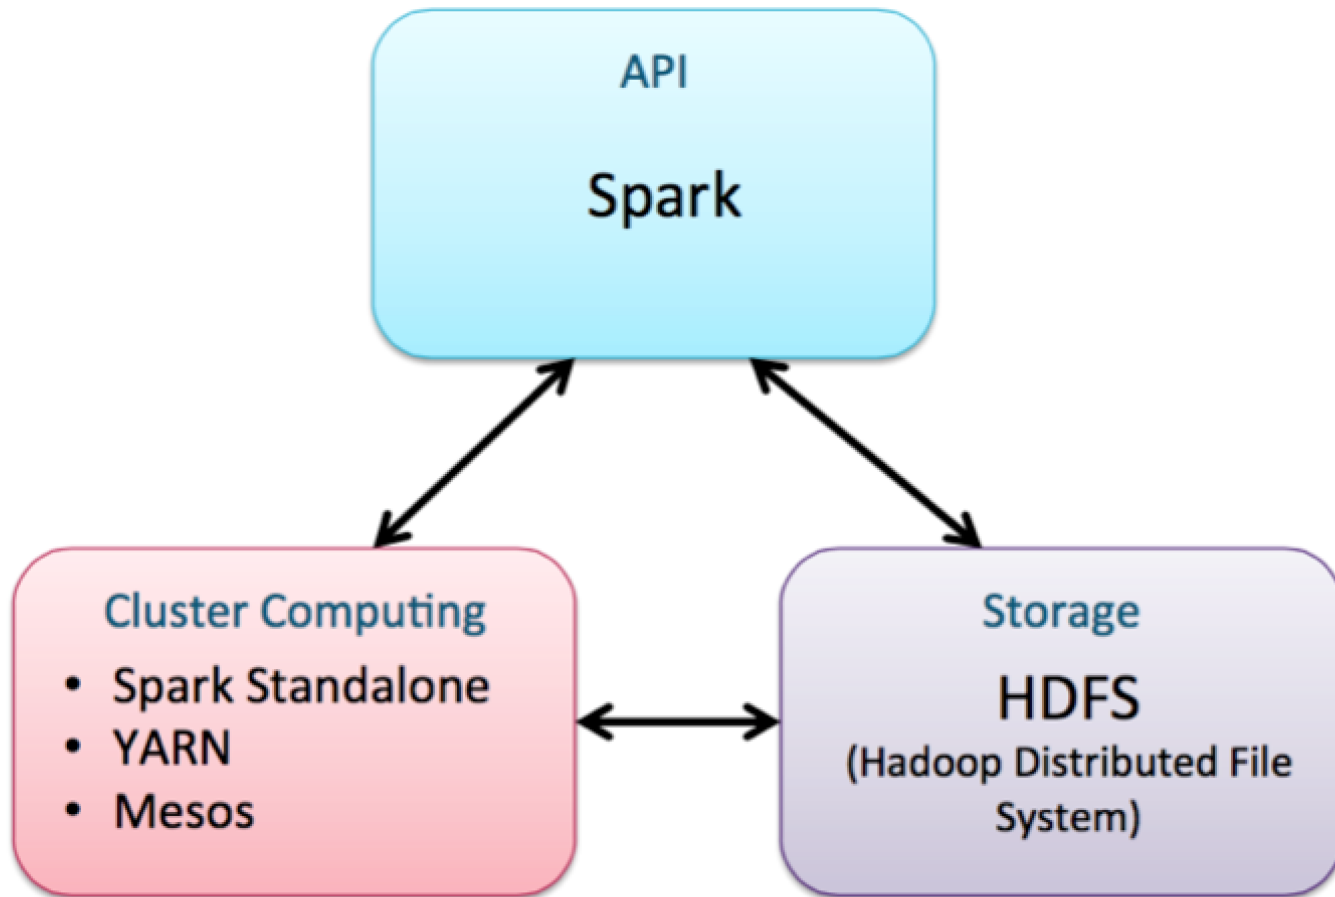
Performance Benchmarks

13



Distributed Computing with Spark

14



Unify Real-Time and Historical Analytics

15

- Spark allows one to write virtually the same batch and streaming codes
 - ▣ Easy to develop and maintain consistency


```
// count words from a file (batch)  
val file = sc.textFile("hdfs://.../pagecounts-*.gz")  
val words = file.flatMap(line => line.split(" "))  
val wordCounts = words.map(x => (x, 1)).reduceByKey(_ + _)  
wordCounts.print()
```

```
// count words from a network stream, every 10s (streaming)  
val ssc = new StreamingContext(args(0), "NetCount", Seconds(10), ..)  
val lines = ssc.socketTextStream("localhost", 3456)  
val words = lines.flatMap(_.split(" "))  
val wordCounts = words.map(x => (x, 1)).reduceByKey(_ + _)  
wordCounts.print()  
ssc.start()
```



Spark v. Hadoop MapReduce

16

- Spark takes the concepts of MapReduce to the next level
 - ▣ Higher-level API = faster, easier development
 - ▣ Low latency = near real-time processing
 - ▣ In-memory data storage = up to 100x performance improvement



```
sc.textFile(file) \  
  .flatMap(lambda s: s.split()) \  
  .map(lambda w: (w,1)) \  
  .reduceByKey(lambda v1,v2: v1+v2) \  
  .saveAsTextFile(output)
```



```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        Job job = new Job();  
        job.setJarByClass(WordCount.class);  
        job.setJobName("Word Count");  
        FileInputFormat.setInputPaths(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
        job.setMapperClass(WordMapper.class);  
        job.setReducerClass(SumReducer.class);  
        job.setMapOutputKeyClass(Text.class);  
        job.setMapOutputValueClass(IntWritable.class);  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
        boolean success = job.waitForCompletion(true);  
        System.exit(success ? 0 : 1);  
    }  
}  
  
public class WordMapper extends Mapper<LongWritable, Text, Text,  
    IntWritable> {  
    public void map(LongWritable key, Text value,  
        Context context) throws IOException, InterruptedException {  
        String line = value.toString();  
        for (String word : line.split("\\W+")) {  
            if (word.length() > 0)  
                context.write(new Text(word), new IntWritable(1));  
        }  
    }  
}  
  
public class SumReducer extends Reducer<Text, IntWritable, Text,  
    IntWritable> {  
    public void reduce(Text key, Iterable<IntWritable>  
        values, Context context) throws IOException, InterruptedException {  
        int wordCount = 0;  
        for (IntWritable value : values) {  
            wordCount += value.get();  
        }  
        context.write(key, new IntWritable(wordCount));  
    }  
}
```


What is Apache Spark

17

- Apache Spark is a fast and general engine for large-scale data processing
- Written in Scala
 - ▣ Functional programming language that runs in a JVM
- Spark Shell
 - ▣ Interactive – for learning or data exploration
 - ▣ Python or Scala
- Spark Application
 - ▣ For large scale data processing
 - ▣ Python, Scala, or Java

RDD (Resilient Distributed Dataset)

18

□ Resilient Distributed Datasets

- Collections of objects spread across a cluster, stored in RAM or on Disk
 - Analogous to HDFS but in memory
 - Still works efficiently on disks
- Resilient – if data in memory is lost, it can be recreated
- Distributed – stored in memory across the cluster
- Dataset – initial data can come from a file or be created programmatically
- RDDs are the fundamental unit of data in Spark
- Most Spark programming consists of performing operations on RDDs
- Automatically rebuilt on failure

RDD Operations

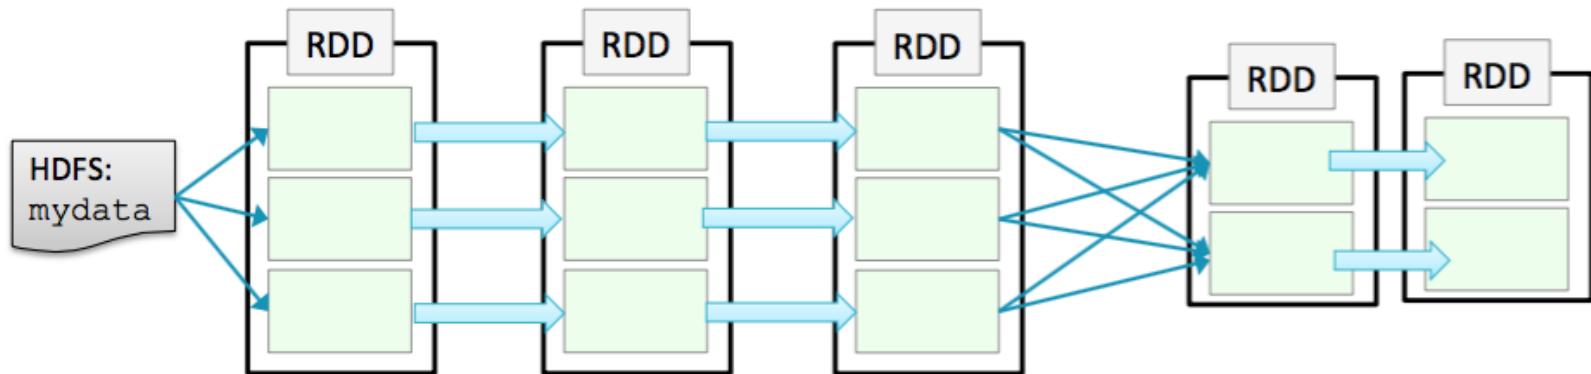


19

□ **Transformations** → create new RDDs

▣ (e.g. map, filter, groupBy)

□ **Actions** → returns value


▣ (e.g. count, collect, save)



RDD Operations: Transformations

20

- Transformations create a new RDD from an existing one
- RDDs are immutable
 - ▣ Data in an RDD is never changed
 - ▣ Transform in sequence to modify the data as needed
- Some common transformations
 - ▣ **map**(function) – creates a new RDD by performing a function on each record in the base RDD
 - ▣ **filter**(function) – creates a new RDD by including or excluding each record in the base RDD according to a boolean function



RDD Operations

21

- | | | |
|------------------|---------------|-------------|
| □ map | □ reduce | sample |
| □ filter | □ count | take |
| □ groupBy | □ fold | first |
| □ sort | □ reduceByKey | partitionBy |
| □ union | □ groupByKey | mapWith |
| □ join | □ cogroup | pipe |
| □ leftOuterJoin | □ cross | save |
| □ rightOuterJoin | □ zip | ... |

Good Explanations:

<https://trongkhoanguyenblog.wordpress.com/2014/11/27/understand-rdd-operations-transformations-and-actions/>

RDD Operations Explained

Loading messages from a log into memory and search for various patterns


```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()

messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
. . .
```

Base RDD

Transformed RDD

Action



Starting Spark Shell

23

- The Spark Shell provides interactive data exploration
- Similar to the “hive” or “pig” command
- Set environment variable in linux
 - ▣ `export SPARK_HOME=/usr/hdp/2.3.2.0-2950/spark/`
 - ▣ `export PATH=/usr/hdp/2.3.2.0-2950/spark/bin:$PATH`

Python Shell: **pyspark**

```
$ pyspark

Welcome to

      _ _ _ _ _
     / / / / /
    / / / / /
   / / / / /
  / / / / /
 / / / / /
/ / / / /
version 1.0.0

Using Python version 2.6.6 (r266:84292, Jan
22 2014 09:42:36)
SparkContext available as sc.

>>>
```

Creating RDDs

Turn a Python collection into an RDD

> sc.parallelize([1, 2, 3])

Load text file from local FS

sc.textFile("file:///root/lab/full_text.txt")

Load text file from HDFS

> sc.textFile("/user/root/shakespeare_100.txt")

> sc.textFile("/user/root/genre/*")

> sc.textFile("hdfs://sandbox.hortonworks.com:8020/user/root/shakespeare_100.txt")

Basic Transformations

```
nums = sc.parallelize([1, 2, 3])  
text = sc.textFile("/user/root/shakespeare_100.txt")
```

Map each element to zero or more others

```
nums.flatMap(lambda x: range(x))  
words = text.map(lambda line: line.split())
```

Pass each element through a function

```
squares = nums.map(lambda x: x*x)  
wordWithCount = words.map(lambda word: (word, 1))
```

Keep elements passing a predicate

```
even = squares.filter(lambda x: x % 2 == 0)
```

Basic Actions

- > `nums = sc.parallelize([1, 2, 3])`
- # Retrieve RDD contents as a local collection
 - > `nums.collect()` # => `[1, 2, 3]`
- # Return first K elements
 - > `nums.take(2)` # => `[1, 2]`
- # Count number of elements
 - > `nums.count()` # => `3`
- # Merge elements with an associative function
 - > `nums.reduce(lambda x, y: x + y)` # => `6`
- # Write elements to a text file in HDFS
 - > `nums.saveAsTextFile("file.txt")`
 - > # To save to local file system
 - > `X = nums.collect()`
 - > Then save using standard write operations

Using spark-submit

27

- Submit some spark commands as a python file

sparkTemplate.py

```
from pyspark import SparkConf, SparkContext
```

```
def main(sc):
```

```
    textFile = sc.textFile("shakespeare_100.txt")
```

```
    words = textFile.flatMap(lambda line: line.split())
```

```
    wordWithCount = words.map(lambda word: (word, 1))
```

```
    wordWithCount.saveAsTextFile("word_count.txt")
```

```
    collectedWords = wordWithCount.collect()
```

```
    print(collectedWords)
```

```
if __name__ == "__main__":
```

```
    conf = SparkConf().setAppName("Testing Spark Commands")
```

```
    sc = SparkContext(conf = conf)
```

```
    main(sc)
```

```
    sc.stop()
```

- Submit job using

- ▣ *spark-submit --master yarn-client --executor-memory 512m --num-executors 3 --executor-cores 1 --driver-memory 512m sparkTemplate.py*

<http://spark.apache.org/docs/latest/submitting-applications.html>

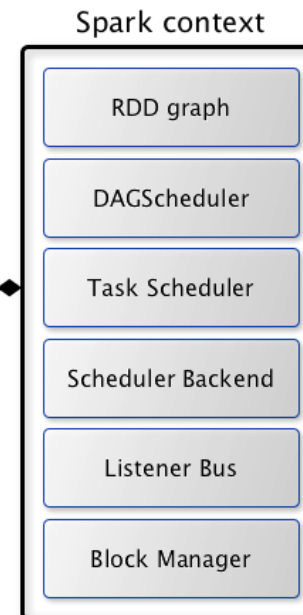
<http://hortonworks.com/hadoop-tutorial/a-lap-around-apache-spark/>

Spark Context

28

- SparkContext (aka Spark context) represents the connection to a Spark execution environment
- You have to create a Spark context before using Spark features and services in your application.
- A Spark context can be used to create RDDs, accumulators and broadcast variables, access Spark services and run jobs.

```
val sc = new SparkContext(master="local[*]",  
    appName="SparkMe App", new SparkConf)  
val lines = sc.textFile(...).cache()  
val c = lines.count()  
println(s"There are $c lines in $fileName")
```



<https://jaceklaskowski.gitbooks.io/mastering-apache-spark/content/spark-sparkcontext.html>

RDD Operations

29

File: purplecow.txt

I've never seen a purple cow.
I never hope to see one;
But I can tell you, anyhow,
I'd rather see than be one.

RDD Operations


30

- `mydata =
sc.textFile("shakespeare_100.txt")`

File: purplecow.txt

I've never seen a purple cow.
I never hope to see one;
But I can tell you, anyhow,
I'd rather see than be one.

RDD: mydata



RDD Operations


31

- `mydata =
sc.textFile("shakespeare_100.txt")`
- `mydata_uc = mydata.map(lambda line:
line.upper())`


File: purplecow.txt

I've never seen a purple cow.
I never hope to see one;
But I can tell you, anyhow,
I'd rather see than be one.

RDD: mydata



RDD: mydata_uc



RDD Operations


32

- `mydata = sc.textFile("shakespeare_100.txt")`
- `mydata_uc = mydata.map(lambda line: line.upper())`
- `mydata_filt = mydata_uc.filter(lambda line: line.startswith('I'))`

File: purplecow.txt

I've never seen a purple cow.
I never hope to see one;
But I can tell you, anyhow,
I'd rather see than be one.


RDD: mydata



RDD: mydata_uc



RDD: mydata_filt



RDD Operations

33

- `mydata =`
`sc.textFile("shakespeare_100.txt")`
- `mydata_uc = mydata.map(lambda line:`
`line.upper())`
- `mydata_filt = mydata_uc.filter(lambda`
`line: line.startswith('I'))`
- `mydata_filt.count()` Action

3

File: purplecow.txt

I've never seen a purple cow.
I never hope to see one;
But I can tell you, anyhow,
I'd rather see than be one.

RDD: `mydata`

I've never seen a purple cow.
I never hope to see one;
But I can tell you, anyhow,
I'd rather see than be one.

RDD: `mydata_uc`

I'VE NEVER SEEN A PURPLE COW.
I NEVER HOPE TO SEE ONE;
BUT I CAN TELL YOU, ANYHOW,
I'D RATHER SEE THAN BE ONE.

RDD: `mydata_filt`

I'VE NEVER SEEN A PURPLE COW.
I NEVER HOPE TO SEE ONE;
I'D RATHER SEE THAN BE ONE.

RDD Map/Reduce Operations

34

- MapReduce in Spark works on Pair RDDs
- Spark's “distributed reduce” transformations operate on RDDs of key-value pairs

Python:

```
pair = (a, b)  
pair[0] # => a  
pair[1] # => b
```

Creating Pair RDDs

35

- The first step in most workflows is to get the data into key/value form
 - ▣ What should the RDD be keyed on?
 - ▣ What is the value?
- Commonly used functions to create Pair RDDs
 - ▣ map
 - ▣ flatMap / flatMapValues
 - ▣ keyBy


Example: A Simple Pair RDD

36

```
➤ text = sc.textFile("full_text.txt") \
    .map(lambda line: line.split('\t')) \
    .map(lambda fields: (fields[0], fields[1]))
```

```
u'USER_79321756\t2010-03-03T04:15:26\t\xdcT: 47.528139,-122.197916\t47.528139\t-122.197916\tRT @USER_2ff4faca: IF SHE DO IT 1 MORE TIME.....IMA KNOCK HER DAMN KOOFIE OFF.....ON MY MOMMA&gt;&gt;haha. #cutthatout'
```

```
u'USER_79321756\t2010-03-03T04:55:32\t\xdcT: 47.528139,-122.197916\t47.528139\t-122.197916\t@USER_77a4822d @USER_2ff4faca okay:) lol. Saying ok to both of yall about to different things!:*'
```



```
(u'USER_79321756', u'2010-03-03T04:15:26')
```

```
(u'USER_79321756', u'2010-03-03T04:55:32')
```


Example: Keying Friend Pairs by Friend ID

37

```
text = sc.textFile("full_text.txt") \
    .keyBy(lambda line: line.split("\t")[0])
#Adding a key as userID
```

```
u'USER_79321756\t2010-03-03T04:15:26\t\tdcT: 47.528139,-122.197916\t47.528139\t-122.197916\tRT @USER_2ff4faca: IF SHE DO IT 1 MORE TIME.....IMA KNOCK HER DAMN KOOFIE OFF.....ON MY MOMMA&gt;&gt;haha. #cutthatout'
```

```
u'USER_79321756\t2010-03-03T04:55:32\t\tdcT: 47.528139,-122.197916\t47.528139\t-122.197916\t@USER_77a4822d @USER_2ff4faca okay:) lol. Saying ok to both of yall about to different things!:*'
```



```
(u'USER_79321756',
u'USER_79321756\t2010-03-03T04:15:26\t\tdcT: 47.528139,-122.197916\t47.528139\t-122.197916\tRT @USER_2ff4faca: IF SHE DO IT 1 MORE TIME.....IMA KNOCK HER DAMN KOOFIE OFF.....ON MY MOMMA&gt;&gt;haha. #cutthatout')
```

```
(u'USER_79321756',
u'USER_79321756\t2010-03-03T04:55:32\t\tdcT: 47.528139,-122.197916\t47.528139\t-122.197916\t@USER_77a4822d @USER_2ff4faca okay:) lol. Saying ok to both of yall about to different things!:*')
```


Example: Pairs with Complex Values

38

```
text = sc.textFile("full_text.txt") \
    .map(lambda line: line.split('\t')) \
    .map(lambda fields: (fields[0], (fields[1], fields[2])))
```

u'USER_79321756\t2010-03-03T04:15:26\t\xdcT: 47.528139,-122.197916\t47.528139\t-122.197916\tRT @USER_2ff4faca: IF SHE DO IT 1 MORE TIME.....IMA KNOCK HER DAMN KOOFIE OFF.....ON MY MOMMA>>haha. #cutthatout'

u'USER_79321756\t2010-03-03T04:55:32\t\xdcT: 47.528139,-122.197916\t47.528139\t-122.197916\t@USER_77a4822d @USER_2ff4faca okay:) lol. Saying ok to both of yall about to different things!:*'



(u'USER_79321756', (u'2010-03-03T04:15:26', u'\xdcT: 47.528139,-122.197916'))

u'USER_79321756', (u'2010-03-03T04:55:32', u'\xdcT: 47.528139,-122.197916'))

Some Key-Value Operations


- > `pets = sc.parallelize([("cat", 1), ("dog", 1), ("cat", 2)])`
- > `pets.reduceByKey(lambda x, y: x + y) # => {(cat, 3), (dog, 1)}`
- > `pets.groupByKey() # => {(cat, [1, 2]), (dog, [1])}`
- > `pets.sortByKey() # => {(cat, 1), (cat, 2), (dog, 1)}`

reduceByKey also automatically implements
combiners on the map side

Spark RDD: WordCount Example

40

the cat sat on the mat
the aardvark sat on the sofa



aardvark	1
cat	1
mat	1
on	2
sat	2
sofa	1
the	4

Spark RDD: WordCount Example

41

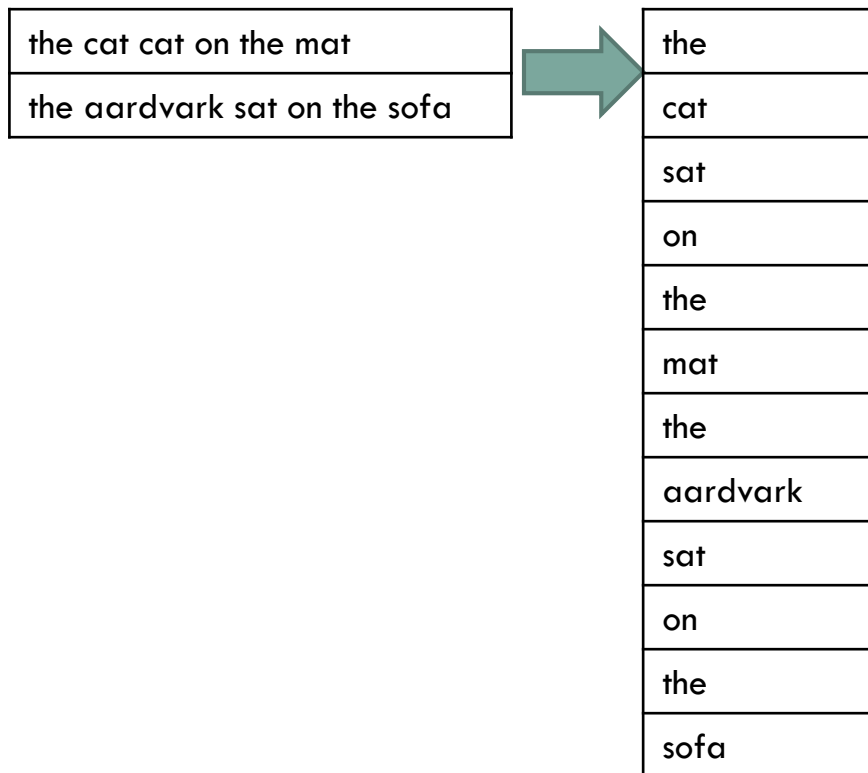
➤ counts =
sc.textFile("shakespeare_100.txt")

the cat cat on the mat
the aardvark sat on the sofa

Spark RDD: WordCount Example

42

- `counts = sc.textFile(text)
 .flatMap(lambda line: line.split())`

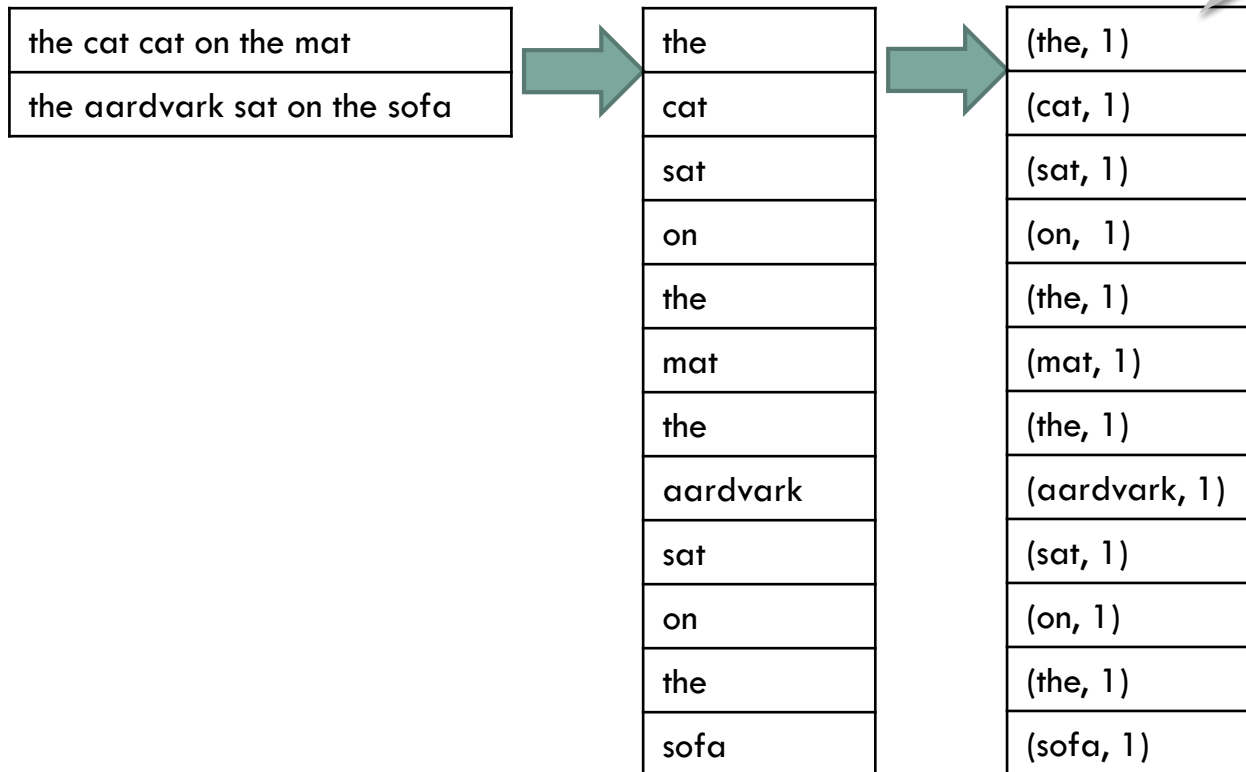


Spark RDD: WordCount Example

43

```
counts = sc.textFile(text)
          .flatMap(lambda line: line.split() )
          .map(lambda word: (word,1) )
```

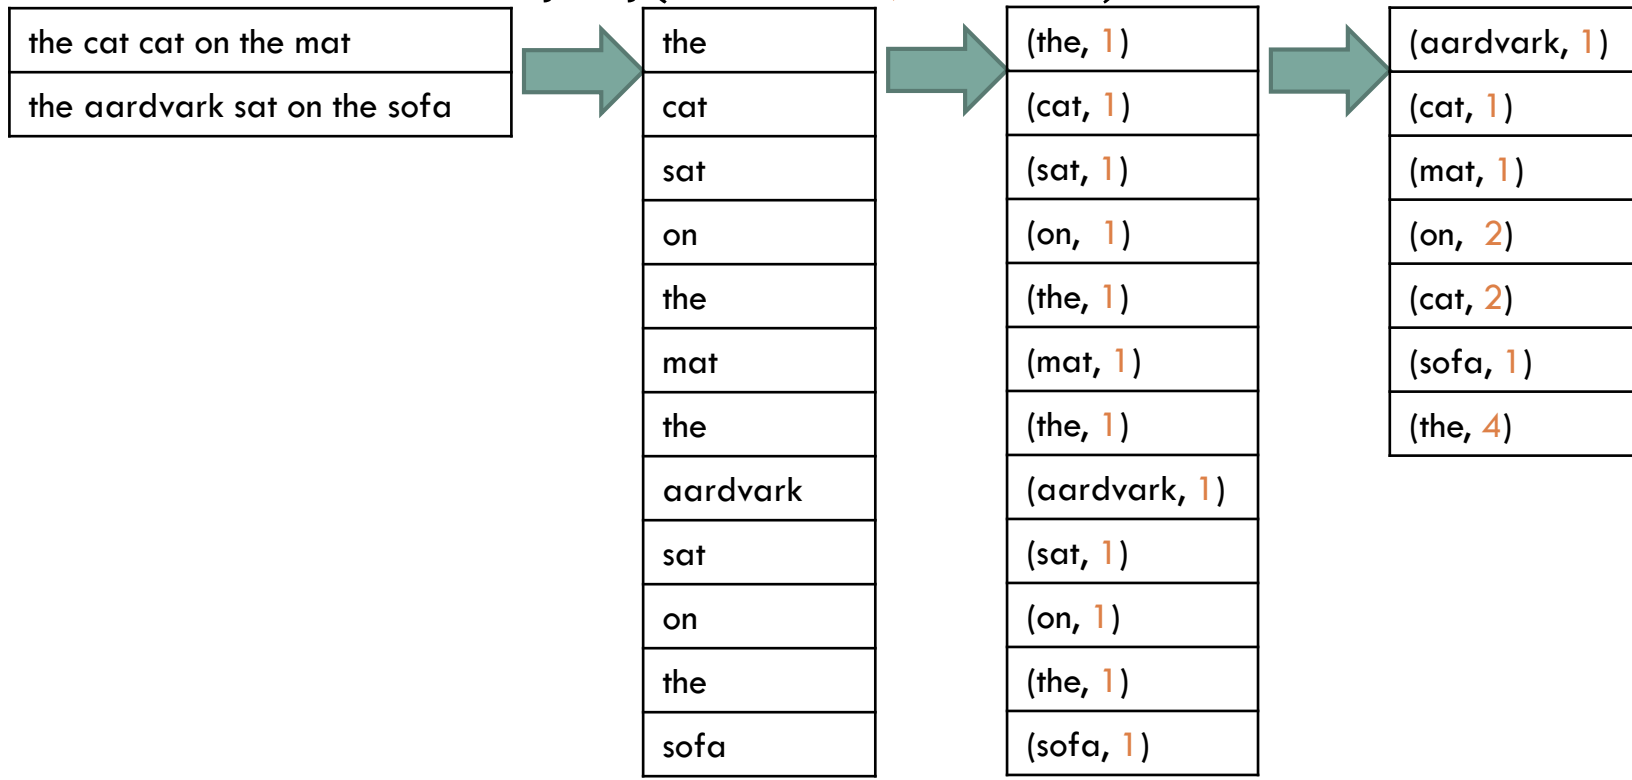
key-value
pairs



Spark RDD: WordCount Example

44

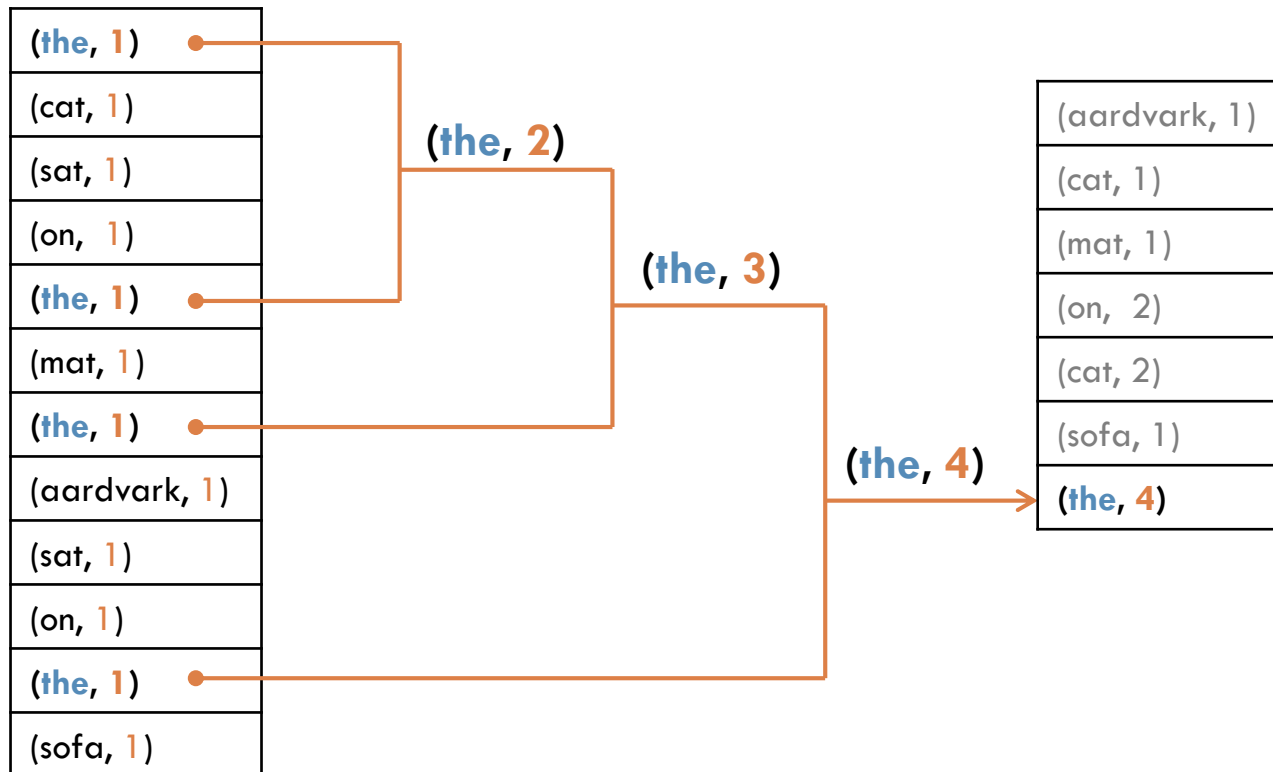
```
counts = sc.textFile(text)
    .flatMap(lambda line: line.split() )
    .map(lambda word: (word,1) )
    .reduceByKey(lambda v1,v2: v1+v2)
```



Spark RDD: WordCount Example

45

- ReduceByKey functions must be
 - ▣ Binary – combines values from two keys
 - ▣ Commutative $\rightarrow x+y = y+x$
 - ▣ Associative $\rightarrow (x+y)+z = x+(y+z)$



Monitor Status of Spark Job

46

- All jobs run on the VM can be monitored using

<http://127.0.0.1:8088/cluster>

The screenshot shows the Hadoop All Applications page. The top section displays the Hadoop logo and the title "All Applications". Below this, there are two main sections: "Cluster Metrics" and "Scheduler Metrics".

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Reb No
13	0	1	12	2	1.95 GB	2.20 GB	0 B	2	8	0	1	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[MEMORY]	<memory:250, vCores:1>	<memory:2250, vCores:8>

Below the scheduler metrics, there is a table of applications. The table has columns for ID, User, Name, Application Type, Queue, StartTime, FinishTime, State, FinalStatus, Running Containers, Progress, Tracking UI, and Blackboard. The first row shows an application with ID "application_1458079610402_0013", User "root", Name "MyApp", Application Type "SPARK", Queue "default", StartTime "Wed Mar 16 20:39:28 -0400 2016", FinishTime "N/A", State "RUNNING", FinalStatus "UNDEFINED", Running Containers "2", Progress "0", Tracking UI "ApplicationMaster", and Blackboard "0".

Click to
get to
the page
below

The screenshot shows the Spark Jobs page. The top section displays the Spark logo and the title "Spark Jobs". Below this, there are two main sections: "Active Jobs" and "Completed Jobs".

Active Jobs (1)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	saveAsTextFile at NativeMethodAccessorImpl.java:-2	2016/03/17 00:40:00	0.6 s	0/2	0/4

Completed Jobs (1)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
0	collectAsMap at /root/deptSum.py:10	2016/03/17 00:39:55	5 s	2/2	4/4

Getting Started with Spark

47

- ❑ Install Spark Standalone on Linux
- ❑ Cloudera/HDP Distributions
 - ❑ http://www.cloudera.com/content/cloudera/en/downloads/quickstart_vms/cdh-5-3-x.html

Spark Documentation

Spark 1.0.0 Documentation: <http://spark.apache.org/docs/latest/>

Spark Programming Guide (Scala,

Python): <http://spark.apache.org/docs/latest/programming-guide.html#overview>

Spark Cassandra Connector - DataStax ([github](#))

Spark HBase - lighting Spark with HBase ([link](#))

Spark MLlib Documentation ([link](#))

Spark GraphX Documentation ([link](#))

(Spark) Spark Cluster Mode Overview ([link](#))

(Spark) Running Spark on EC2 ([link](#))

(Cloudera) Pig is Flying: Apache Pig on Apache Spark ([link](#))